


I Want To Go Faster!

A Beginner's Guide to Indexing



@bertwagner



bertwagner.com



youtube.com/bertwagner



bert@bertwagner.com

Bert Wagner



Slides available here!



@bertwagner



bertwagner.com



youtube.com/bertwagner



bert@bertwagner.com

Why Indexes?

- Biggest bang for the buck
 - Can potentially fix many queries at once
- Positive downstream side effects
 - Can help:
 - Reduce blocking
 - Prevent deadlocking
 - Improve caching





Why **NOT** Indexes?

- They take up space
- Every index adds overhead on insert, updates, deletes
- Maintenance
 - Fragmentation – external and internal
 - Ownership
 - Maintenance windows introduce blocking, downtime, coordination




When you need to index

- You have many different queries running slowly against a table (or multiple tables)
- You have one query with lots of executions is slow (refactor first)
- You see Microsoft's green index recommendations in execution plans



- Seeing lots of scans in execution plans (not always bad)





Index Internals



@bertwagner



bertwagner.com



youtube.com/bertwagner



bert@bertwagner.com

Heap



Blue Jay


Color: Blue

Size: 4 inches

Description:
Territorial, loud
and obnoxious
call.

Habitat:
Trees, bushes,
feeders in suburban backyards.

Migration:
Stays year round.



- Table with unordered rows of data
- Finding a particular row always requires a scan



@bertwagner



bertwagner.com



youtube.com/bertwagner



bert@bertwagner.com


Heap



```
CREATE TABLE dbo.FieldGuide  
(  
    BirdName nvarchar(100),  
    Color nvarchar(20),  
    Size tinyint,  
    Description nvarchar(1000),  
    Habitat nvarchar(1000),  
    Migration nvarchar(1000)  
)
```



Clustered Index



Ordered
alphabetically
by bird name.

- Default table order is defined – rows stored sorted in that order.
- Finding a particular row can now sometimes use a seek.

Blue Jay


Color: Blue

Size: 4 inches

Description:
Territorial, loud and obnoxious call.

Habitat:
Trees, bushes, feeders in suburban backyards.

Migration:
Stays year round.



Cardinal


Color: Red

Size: 6 inches

Description:
Red with a mohawk on its head.

Habitat:
Trees, bushes, feeders in suburban backyards.

Migration:
Stays year round.



Sparrow

Color: Brown

Size: 2 inches

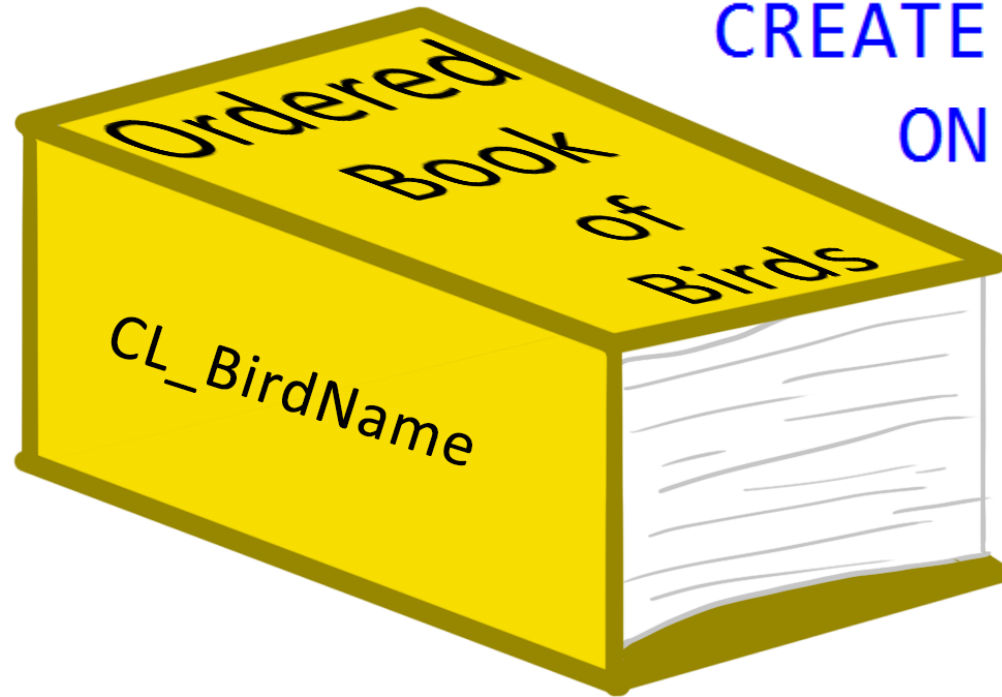
Description:
Little brown bird.

Habitat:
Tall bushes, thick growth.

Migration:
Stays year round.



Clustered Index




```
CREATE CLUSTERED INDEX CL_BirdName  
ON dbo.FieldGuide (BirdName);
```

Our index “key” – what
order the data is stored in.



Nonclustered Index



Ordered
alphabetically
by bird name...
but with fewer
columns

- Subset of columns
- Stored in a different order

Blue Jay	
Color: Blue	
Size: 4 inches	
Cardinal	
Color: Red	
Size: 6 inches	
Crow	
Color: Black	
Size: 12 inches	


Flamingo	
Color: Pink	
Size: 36 inches	
Goldfinch	
Color: Yellow	
Size: 3 inches	
Oriole	
Color: Orange	
Size: 4 inches	

Seagull	
Color: White	
Size: 14 inches	
Sparrow	
Color: Brown	
Size: 2 inches	
Titmouse	
Color: Grey	
Size: 3 inches	



Nonclustered Index


```
CREATE NONCLUSTERED INDEX IX_BirdName_Includes  
ON dbo.FieldGuide (BirdName)  
INCLUDE (Color, Size);
```



Our “included” columns – these are the columns of data that get copied into our index.



Nonclustered Index Part 2



- Same subset of columns as previous nonclustered index
- Sorted in different order, so better for certain queries (and worse for others)

Crow Color: Black Size: 12 inches	
Blue Jay Color: Blue Size: 4 inches	
Sparrow Color: Brown Size: 2 inches	


Titmouse Color: Grey Size: 3 inches	
Oriole Color: Orange Size: 4 inches	
Flamingo Color: Pink Size: 36 inches	

Cardinal Color: Red Size: 6 inches	
Seagull Color: White Size: 14 inches	
Goldfinch Color: Yellow Size: 3 inches	



Nonclustered Index Part 2

```
CREATE NONCLUSTERED INDEX IX_Color_Includes  
ON dbo.FieldGuide (Color)  
INCLUDE (BirdName, Size);
```



Cardinality

	Name	Color	Size
1.	Blue Jay	Blue	Medium
2.	Cardinal	Red	Medium
3.	Crow	Black	Large
4.	Flamingo	Pink	Extra Large
5.	Oriole	Orange	Medium
6.	Sparrow	Brown	Small

```
SELECT
    Name, Color, Size
FROM
    dbo.Birds
WHERE
    Color = 'Red'
    AND Size = 'Medium'
```

```
CREATE INDEX IX_Size
ON dbo.Birds (Size,Color)
INCLUDE (Name)
```

```
CREATE INDEX IX_Color
ON dbo.Birds (Color,Size)
INCLUDE (Name)
```

Rows 1-20	Extra Small
Rows 21-40	Small
Rows 41-60	Medium
Rows 61-80	Large
Rows 81-100	Extra Large

Rows 1-5	Amber
Rows 6-10	Black
Rows 11-15	Blue
Rows 16-20	Brown
Rows 21-25	Fuscia
Rows 26-30	Gold
Rows 31-35	Green
Rows 36-40	Grey
Rows 41-45	Hazel
Rows 46-50	Light Blue
Rows 51-55	Mauve
Rows 56-60	Orange
Rows 61-65	Pink
Rows 66-70	Purple
Rows 71-75	Red
Rows 76-80	Scarlet
Rows 81-85	Silver
Rows 86-90	Tangerine
Rows 91-95	White
Rows 96-100	Yellow



Summary

Indexes solve two big problems:

1. Sort Your Data
 - JOIN, WHERE, GROUP BY, etc... can utilize this pre-sorted data.
2. Allow for Higher Data Density
 - So there are fewer data pages to read



Any questions so far?



@bertwagner



bertwagner.com



youtube.com/bertwagner



bert@bertwagner.com

QUIZ TIME!

- Does key column order matter?
 - Imagine a book of birds sorted on bird name, will it help me find unknown birds?
 - No – it's better to sort on some other attribute first, such as color.
 - Imagine my book has 100 birds – evenly distributed across 20 distinct colors and 5 different sizes. If I want to find brown birds that are 3" tall, do I first sort on color or on size?
 - Color – it's more selective. Filtering on a single color will leave me with 5 possible bird choices. Filtering on size first would leave me with 20 possible bird choices.
- Does include column order matter?
 - Include columns are primarily there to prevent your queries from having to go to a different index to get additional columns. Order doesn't matter here.



Common Indexing Recommendations

- These are suggestions, not absolutes!
- They are good to get you going, but they won't be the best option in every scenario (especially complex situations).
- But they are good starting points.



Heap

- Benefits:
 - If no clustered index, it means data can be written quickly wherever there is free space. *Possibly* useful in ETL staging.
- Disadvantages:
 - A well-defined clustered index will almost always beat out a heap.
 - Every operation that is not an insert will be as fast or significantly faster on a clustered index than on a heap.
- Recommendation:
 - Don't use them.



Heap

```
CREATE TABLE dbo.BirdSpecies  
(  
    Id bigint IDENTITY,  
    CommonName nvarchar(100),  
    PrimaryColor nvarchar(30),  
    SecondaryColor nvarchar(30),  
    HeightInches tinyint  
);
```



Clustered Index

- Benefits:
 - Better than a heap 99% of the time.
 - Can make lookups quicker.
- Disadvantages:
 - A poorly defined clustered index could take up extra space and be inefficient.
- Other Notes:
 - You get one per table – choose wisely!
 - The clustering key gets copied to every nonclustered index page...don't make it wide (either # of columns or byte size of columns)



Clustered Index

- Recommended keys for a clustered index:
 - **Int/bigint IDENTITY** – will guarantee uniqueness, ever-increasing.
 - If I don't know what to do, this is my default.
 - **Datetime2** – ever increasing, may not be unique so might want to add a second IDENTITY column as part of the key.
 - This is particularly useful if you are always going to be querying on that datetime2 field.
 - **Sequential GUIDs** – ever increasing, unique across systems
 - Only would use this option **if** need a unique value across tables, servers.



Clustered Index - Identity

```
CREATE TABLE dbo.FieldObservations
(
    Id bigint IDENTITY,
    BirdSpeciesId int,
    DateSeen datetime2,
    LocationLatLong geography,
    ObserverName nvarchar(100),
    CONSTRAINT PK_FieldObservationsId PRIMARY KEY CLUSTERED (Id)
);
```



A Quick Note about Primary Keys

- PKs != Clustered Indexes
- PKs are a constraint that indicate column combination makes a row unique
- You can have PK as a nonclustered index. Or a clustered index without a PK. They're independent of each other.



Clustered Index – Datetime2 + identity

```
CREATE TABLE dbo.FieldObservations
(
    DateSeen datetime2,
    Id bigint IDENTITY,
    BirdSpeciesId int,
    LocationLatLong geography,
    ObserverName nvarchar(100),
    CONSTRAINT PK_DateSeen_Id PRIMARY KEY CLUSTERED (DateSeen, Id)
);
```



Clustered Index – Sequential GUID

```
CREATE TABLE dbo.FieldObservations
(
    Id uniqueidentifier CONSTRAINT DF_FieldObservationId DEFAULT NEWSEQUENTIALID(),
    BirdSpeciesId int,
    DateSeen datetime2,
    LocationLatLong geography,
    ObserverName nvarchar(100),
    CONSTRAINT PK_FieldObservationId PRIMARY KEY CLUSTERED (Id)
);
```



Nonclustered Index

- Benefits:
 - Can store data sorted in a different order.
 - Can store data with fewer columns (greater density)
- Disadvantages:
 - Use extra space.
 - Need to be modified on every insert/update/delete.
 - Need to be maintained for fragmentation.



Nonclustered Index

- Recommended keys for a nonclustered index:
 - **Foreign keys** – fields used to join on (regardless if they have a FK constraint or not)
 - **WHERE predicates** – if a query is filtering on a subset of rows
 - **GROUP BYs** – if using GROUP BYs or window functions, use your grouping/partitioning columns as keys



Nonclustered Index

-- Who observed birds on March 1st?

```
SELECT
    fo.ObserverName, fo.DateSeen
FROM
    dbo.FieldObservations fo
WHERE
    fo.DateSeen >= '2018-03-01'
    AND fo.DateSeen < '2018-03-02'
```

```
CREATE TABLE dbo.FieldObservations
(
    Id bigint IDENTITY,
    BirdSpeciesId int,
    DateSeen datetime2,
    LocationLatLong geography,
    ObserverName nvarchar(100),
    CONSTRAINT PK_FieldObservationsId PRIMARY KEY CLUSTERED (Id)
);
```

```
CREATE NONCLUSTERED INDEX IX_DateSeen_Includes
ON dbo.FieldObservations (DateSeen) INCLUDE (ObserverName)
```



Nonclustered Index

-- Who observed birds on March 1st?

```
SELECT
    fo.ObserverName, fo.DateSeen
FROM
    dbo.FieldObservations fo
WHERE
    fo.DateSeen >= '2018-03-01'
    AND fo.DateSeen < '2018-03-02'
    AND fo.ObserverName = 'Bert'
```

```
CREATE TABLE dbo.FieldObservations
(
    Id bigint IDENTITY,
    BirdSpeciesId int,
    DateSeen datetime2,
    LocationLatLong geography,
    ObserverName nvarchar(100),
    CONSTRAINT PK_FieldObservationsId PRIMARY KEY CLUSTERED (Id)
);
```

```
CREATE NONCLUSTERED INDEX IX_DateSeen_ObserverName
ON dbo.FieldObservations (DateSeen, ObserverName)
```

This assumes DateSeen is more selective (has more unique values) than ObserverName



Nonclustered Index

```
-- How many sightings per observer?  
SELECT bs.ObserverName, COUNT(bs.Id) AS ObserverCount  
FROM  
    dbo.FieldObservations bs  
GROUP BY  
    bs.ObserverName  
  
CREATE NONCLUSTERED INDEX IX_ObserverName  
ON dbo.FieldObservations (ObserverName)
```

(A one column index is probably not best practice. Perhaps we can modify this index to be used by other queries as well...)



@bertwagner



bertwagner.com



youtube.com/bertwagner



bert@bertwagner.com

Nonclustered Index

```
-- How many times have I seen each bird?
SELECT  fo.ObserverName,
        COUNT(DISTINCT fo.BirdSpeciesId) AS BirdCount,
        bs.CommonName
FROM
    dbo.FieldObservations fo
    INNER JOIN dbo.BirdSpecies bs
        ON fo.BirdSpeciesId = bs.Id
WHERE
    fo.ObserverName = 'Bert'
GROUP BY
    fo.ObserverName,
    bs.CommonName
```

```
CREATE TABLE dbo.BirdSpecies
(
    Id bigint IDENTITY,
    CommonName nvarchar(200),
    ScientificName nvarchar(200),
    TaxonRank nvarchar(20),
    Kingdom nvarchar(20),
    Phylum nvarchar(20),
    Class nvarchar(20),
    [Order] nvarchar(20),
    Family nvarchar(50),
    Genus nvarchar(50),
    CONSTRAINT PK_BirdSpeciesId PRIMARY KEY CLUSTERED (Id)
);
GO
```

Definitely:

```
CREATE NONCLUSTERED INDEX IX_ObserverName_BirdSpeciesId
ON dbo.FieldObservations (ObserverName, BirdSpeciesId)
```

Maybe:

```
CREATE NONCLUSTERED INDEX IX_BirdSpeciesId_Includes
ON dbo.BirdSpecies (Id) INCLUDE (CommonName)
```



Nonclustered Index

-- We want a query to search all of our columns

```
DECLARE @SearchValue nvarchar(20);
```

```
SELECT
```

```
    bs.Id, bs.CommonName
```

```
FROM
```

```
    dbo.BirdSpecies bs
```

```
WHERE
```

```
    bs.ScientificName = @SearchValue
```

```
    OR bs.Kingdom = @SearchValue
```

```
    OR bs.Phylum = @SearchValue
```

```
    OR bs.Class = @SearchValue
```

```
    OR bs.[Order] = @SearchValue
```

```
    OR bs.Family = @SearchValue
```

```
    OR bs.[Genus] = @SearchValue
```

```
CREATE TABLE dbo.BirdSpecies
```

```
(
```

```
    Id bigint IDENTITY,
```

```
    CommonName nvarchar(200),
```

```
    ScientificName nvarchar(200),
```

```
    TaxonRank nvarchar(20),
```

```
    Kingdom nvarchar(20),
```

```
    Phylum nvarchar(20),
```

```
    Class nvarchar(20),
```

```
    [Order] nvarchar(20),
```

```
    Family nvarchar(50),
```

```
    Genus nvarchar(50),
```

```
    CONSTRAINT PK_BirdSpeciesId PRIMARY KEY CLUSTERED (Id)
```

```
);
```

```
GO
```

One query per column? Maybe. But probably too painful.

```
CREATE NONCLUSTERED INDEX IX_ScientificName_Includes
```

```
    ON dbo.BirdSpecies (ScientificName) INCLUDE (CommonName);
```

```
CREATE NONCLUSTERED INDEX IX_Kingdom_Includes
```

```
    ON dbo.BirdSpecies (Kingdom) INCLUDE (CommonName);
```

```
CREATE NONCLUSTERED INDEX IX_Phylum_Includes
```

```
    ON dbo.BirdSpecies (Phylum) INCLUDE (CommonName);
```



Nonclustered Index

```
-- We want a query to search all of our columns
-- Might have to unpivot our data first
SELECT
    bs.Id, bs.CommonName
FROM
    dbo.BirdSpecies bs
WHERE
    bs.CategoryValue = @SearchValue
```

```
CREATE TABLE dbo.BirdSpecies
(
    Id bigint IDENTITY,
    CommonName nvarchar(200),
    CategoryName nvarchar(30),
    CategoryValue nvarchar(200)
    CONSTRAINT PK_BirdSpeciesId PRIMARY KEY CLUSTERED (Id)
);
```

```
CREATE NONCLUSTERED INDEX IX_CategoryValue_Include
    ON dbo.BirdSpecies (CategoryValue) INCLUDE (CommonName)
```

Not every problem is an index problem.

Might be better to restructure your table instead.



Thank you!

SQL WITH BERT

New Episodes Every
Tuesday



[@bertwagner](https://twitter.com/bertwagner)



bertwagner.com



youtube.com/bertwagner



bert@bertwagner.com

← New posts and videos
every Tuesday!



@bertwagner



bertwagner.com



youtube.com/bertwagner



bert@bertwagner.com

Appendix - When do indexes hurt

- Hotspots - last page is constantly being locked by inserts - this can slow things down. This may be a time where a clustered index may not want to be on something incremental like identity or date. GUIDs actually solve some performance issues here *shudder*
- If you have a query that is slow, but runs once per month for reporting, don't index it. Either deal with the slow tradeoff, or do something like run a job to create that index then remove it after the job runs.
- Insert/Update/Delete - every time this happens, every index must be modified
- page splits
- fill factor
- fragmentation (even ssds, fragmentation stinks bc still reading more pages into buffer pool)
- maintenance - operational burden. someone needs to run maintenance scripts. index maintenance scripts can have lots of problems with blocking. the more indexes you have to manage, the more possibly you have with trouble arising.



Appendix - Other types of indexes

- Filtered, computed column index (storing preparsed data or fixing conversion issues without changing queries), Columnstore (analytical queries), spatial, hierarchical, xml, etc..
- Link to resources

